

# TensorFlow\_2.x\_YOLOv3

2023.02.08

# TensorFlow\_2.x\_YOLOv3

```
[ ] 1 cd /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master  
  
/content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master
```

```
▶ 1 !python mnist/make_data.py
```

```
↳ Extracting all train files now...  
Done!  
Extracting all test files now...  
Done!
```

```
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000001.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000002.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000003.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000004.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000005.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000006.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000007.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000008.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000009.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000010.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000011.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000012.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000013.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000014.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000015.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000016.jpg  
=> /content/drive/MyDrive/TensorFlow-2.x-YOLOv3-master/mnist/mnist_train/000017.jpg
```



# TensorFlow\_2.x\_YOLOv3

```
1 #Training
2 !python train.py

epoch:29 step: 213/250, lr:0.000001, giou_loss: 2.72, conf_loss: 0.35, prob_loss: 2.83, total_loss: 5.90
epoch:29 step: 214/250, lr:0.000001, giou_loss: 2.42, conf_loss: 0.31, prob_loss: 4.77, total_loss: 7.50
epoch:29 step: 215/250, lr:0.000001, giou_loss: 1.89, conf_loss: 0.40, prob_loss: 2.12, total_loss: 4.41
epoch:29 step: 216/250, lr:0.000001, giou_loss: 2.29, conf_loss: 0.33, prob_loss: 1.89, total_loss: 4.51
epoch:29 step: 217/250, lr:0.000001, giou_loss: 1.98, conf_loss: 0.34, prob_loss: 2.06, total_loss: 4.37
epoch:29 step: 218/250, lr:0.000001, giou_loss: 3.12, conf_loss: 0.52, prob_loss: 2.73, total_loss: 6.36
epoch:29 step: 219/250, lr:0.000001, giou_loss: 2.20, conf_loss: 0.46, prob_loss: 2.03, total_loss: 4.69
epoch:29 step: 220/250, lr:0.000001, giou_loss: 2.55, conf_loss: 0.64, prob_loss: 3.08, total_loss: 6.27
epoch:29 step: 221/250, lr:0.000001, giou_loss: 2.16, conf_loss: 0.29, prob_loss: 2.03, total_loss: 4.49
epoch:29 step: 222/250, lr:0.000001, giou_loss: 3.12, conf_loss: 0.33, prob_loss: 2.75, total_loss: 6.21
epoch:29 step: 223/250, lr:0.000001, giou_loss: 2.29, conf_loss: 0.31, prob_loss: 1.91, total_loss: 4.51
epoch:29 step: 224/250, lr:0.000001, giou_loss: 2.82, conf_loss: 0.43, prob_loss: 2.54, total_loss: 5.79
epoch:29 step: 225/250, lr:0.000001, giou_loss: 2.26, conf_loss: 0.36, prob_loss: 2.43, total_loss: 5.05
epoch:29 step: 226/250, lr:0.000001, giou_loss: 2.41, conf_loss: 0.53, prob_loss: 1.97, total_loss: 4.90
epoch:29 step: 227/250, lr:0.000001, giou_loss: 4.54, conf_loss: 3.01, prob_loss: 5.30, total_loss: 12.86
epoch:29 step: 228/250, lr:0.000001, giou_loss: 1.84, conf_loss: 0.42, prob_loss: 2.03, total_loss: 4.30
epoch:29 step: 229/250, lr:0.000001, giou_loss: 1.54, conf_loss: 0.29, prob_loss: 1.38, total_loss: 3.21
epoch:29 step: 230/250, lr:0.000001, giou_loss: 1.46, conf_loss: 0.33, prob_loss: 1.53, total_loss: 3.33
epoch:29 step: 231/250, lr:0.000001, giou_loss: 3.10, conf_loss: 1.85, prob_loss: 3.61, total_loss: 8.56
epoch:29 step: 232/250, lr:0.000001, giou_loss: 2.92, conf_loss: 0.41, prob_loss: 2.94, total_loss: 6.27
epoch:29 step: 233/250, lr:0.000001, giou_loss: 1.83, conf_loss: 0.30, prob_loss: 2.02, total_loss: 4.15
epoch:29 step: 234/250, lr:0.000001, giou_loss: 0.87, conf_loss: 0.27, prob_loss: 0.86, total_loss: 2.00
epoch:29 step: 235/250, lr:0.000001, giou_loss: 3.57, conf_loss: 0.64, prob_loss: 3.63, total_loss: 7.83
epoch:29 step: 236/250, lr:0.000001, giou_loss: 3.74, conf_loss: 2.48, prob_loss: 4.67, total_loss: 10.90
epoch:29 step: 237/250, lr:0.000001, giou_loss: 3.19, conf_loss: 1.91, prob_loss: 3.52, total_loss: 8.63
epoch:29 step: 238/250, lr:0.000001, giou_loss: 2.55, conf_loss: 0.51, prob_loss: 2.92, total_loss: 5.98
epoch:29 step: 239/250, lr:0.000001, giou_loss: 2.82, conf_loss: 0.38, prob_loss: 2.87, total_loss: 6.06
epoch:29 step: 240/250, lr:0.000001, giou_loss: 2.53, conf_loss: 0.48, prob_loss: 2.43, total_loss: 5.43
epoch:29 step: 241/250, lr:0.000001, giou_loss: 2.22, conf_loss: 0.37, prob_loss: 1.78, total_loss: 4.36
epoch:29 step: 242/250, lr:0.000001, giou_loss: 2.44, conf_loss: 0.39, prob_loss: 2.30, total_loss: 5.13
epoch:29 step: 243/250, lr:0.000001, giou_loss: 2.23, conf_loss: 0.47, prob_loss: 1.93, total_loss: 4.63
epoch:29 step: 244/250, lr:0.000001, giou_loss: 2.06, conf_loss: 0.32, prob_loss: 2.08, total_loss: 4.46
epoch:29 step: 245/250, lr:0.000001, giou_loss: 2.39, conf_loss: 0.30, prob_loss: 2.44, total_loss: 5.13
epoch:29 step: 246/250, lr:0.000001, giou_loss: 2.38, conf_loss: 0.58, prob_loss: 2.42, total_loss: 5.38
epoch:29 step: 247/250, lr:0.000001, giou_loss: 2.39, conf_loss: 0.36, prob_loss: 2.67, total_loss: 5.42
epoch:29 step: 248/250, lr:0.000001, giou_loss: 3.68, conf_loss: 1.29, prob_loss: 3.21, total_loss: 8.18
epoch:29 step: 249/250, lr:0.000001, giou_loss: 2.50, conf_loss: 0.35, prob_loss: 2.45, total_loss: 5.30
epoch:29 step: 0/250, lr:0.000001, giou_loss: 2.11, conf_loss: 0.31, prob_loss: 1.80, total_loss: 4.22
epoch:29 step: 1/250, lr:0.000001, giou_loss: 3.21, conf_loss: 0.38, prob_loss: 2.94, total_loss: 6.52

giou_val_loss: 2.48, conf_val_loss: 0.63, prob_val_loss: 4.27, total_val_loss: 7.38
```

```
configs.py ×
29 YOLO_ANCHORS = [[12, 16], [19, 36], [40, 28]],
30 [[36, 75], [76, 55], [72, 146]],
31 [[142, 110], [192, 243], [459, 401]]]
32 if YOLO_TYPE == "yolov3":
33     YOLO_ANCHORS = [[10, 13], [16, 30], [33, 23]],
34 [[30, 61], [62, 45], [59, 119]],
35 [[116, 90], [156, 198], [373, 326]]]
36 # Train options
37 TRAIN_YOLO_TINY = False
38 TRAIN_SAVE_BEST_ONLY = True # saves only best model according validation
39 TRAIN_SAVE_CHECKPOINT = False # saves all best validated checkpoints in t
40 TRAIN_CLASSES = "mnist/mnist.names"
41 TRAIN_ANNOT_PATH = "mnist/mnist_train.txt"
42 TRAIN_LOGDIR = "log"
43 TRAIN_CHECKPOINTS_FOLDER = "checkpoints"
44 TRAIN_MODEL_NAME = f"{YOLO_TYPE}_custom"
45 TRAIN_LOAD_IMAGES_TO_RAM = True # With True faster training, but need more R
46 TRAIN_BATCH_SIZE = 4
47 TRAIN_INPUT_SIZE = 416
48 TRAIN_DATA_AUG = True
49 TRAIN_TRANSFER = True
50 TRAIN_FROM_CHECKPOINT = False # "checkpoints/yolov3_custom"
51 TRAIN_LR_INIT = 1e-4
52 TRAIN_LR_END = 1e-6
53 TRAIN_WARMUP_EPOCHS = 2
54 TRAIN_EPOCHS = 30
```



# TensorFlow\_2.x\_YOLOv3

- train.py  
⇒ `main()` v

```
trainset = Dataset('train')
testset = Dataset('test')

steps_per_epoch = len(trainset)
global_steps = tf.Variable(1, trainable=False, dtype=tf.int64)
warmup_steps = TRAIN_WARMUP_EPOCHS * steps_per_epoch
total_steps = TRAIN_EPOCHS * steps_per_epoch

if TRAIN_TRANSFER:
    Darknet = Create_Yolo(input_size=YOLO_INPUT_SIZE, CLASSES=YOLO_COCO_CLASSES)
    load_yolo_weights(Darknet, Darknet_weights) # use darknet weights

yolo = Create_Yolo(input_size=YOLO_INPUT_SIZE, training=True, CLASSES=TRAIN_CLASSES)

if TRAIN_FROM_CHECKPOINT:
    try:
        yolo.load_weights(f"./checkpoints/{TRAIN_MODEL_NAME}")
    except ValueError:
        print("Shapes are incompatible, transferring Darknet weights")
        TRAIN_FROM_CHECKPOINT = False

if TRAIN_TRANSFER and not TRAIN_FROM_CHECKPOINT:
    for i, l in enumerate(Darknet.layers):
        layer_weights = l.get_weights()
        if layer_weights != []:
            try:
                yolo.layers[i].set_weights(layer_weights)
            except:
                print("skipping", yolo.layers[i].name)

optimizer = tf.keras.optimizers.Adam()
```

# TensorFlow\_2.x\_YOLOv3

- train.py
  - ⇒ main()
  - ⇒ train\_step() v

```
def train_step(image_data, target):
    with tf.GradientTape() as tape:
        pred_result = yolo(image_data, training=True)
        giou_loss=conf_loss=prob_loss=0

        # optimizing process
        grid = 3 if not TRAIN_YOLO_TINY else 2
        for i in range(grid):
            conv, pred = pred_result[i*2], pred_result[i*2+1]
            loss_items = compute_loss(pred, conv, *target[i], i, CLASSES=TRAIN_CLASSES)
            giou_loss += loss_items[0]
            conf_loss += loss_items[1]
            prob_loss += loss_items[2]

        total_loss = giou_loss + conf_loss + prob_loss

        gradients = tape.gradient(total_loss, yolo.trainable_variables)
        optimizer.apply_gradients(zip(gradients, yolo.trainable_variables))

        # update learning rate
        # about warmup: https://arxiv.org/pdf/1812.01187.pdf&usg=ALkJrhglKOPDjNt6SHGbphTHyMcT0cuMJg
        global_steps.assign_add(1)
        if global_steps < warmup_steps:# and not TRAIN_TRANSFER:
            lr = global_steps / warmup_steps * TRAIN_LR_INIT
        else:
            lr = TRAIN_LR_END + 0.5 * (TRAIN_LR_INIT - TRAIN_LR_END)*((1 + tf.cos((global_steps - warmup_steps) / (total_steps - warmup_steps) * np.pi)))
        optimizer.lr.assign(lr.numpy())

        # writing summary data
        with writer.as_default():
            tf.summary.scalar("lr", optimizer.lr, step=global_steps)
            tf.summary.scalar("loss/total_loss", total_loss, step=global_steps)
            tf.summary.scalar("loss/giou_loss", giou_loss, step=global_steps)
            tf.summary.scalar("loss/conf_loss", conf_loss, step=global_steps)
            tf.summary.scalar("loss/prob_loss", prob_loss, step=global_steps)
        writer.flush()

    return global_steps.numpy(), optimizer.lr.numpy(), giou_loss.numpy(), conf_loss.numpy(), prob_loss.numpy(), total_loss.numpy()
```

# TensorFlow\_2.x\_YOLOv3

- train.py  
⇒ main()  
⇒ train\_step() v

```
def validate_step(image_data, target):  
    with tf.GradientTape() as tape:  
        pred_result = yolo(image_data, training=False)  
        giou_loss=conf_loss=prob_loss=0  
  
        # optimizing process  
        grid = 3 if not TRAIN_YOLO_TINY else 2  
        for i in range(grid):  
            conv, pred = pred_result[i*2], pred_result[i*2+1]  
            loss_items = compute_loss(pred, conv, *target[i], i, CLASSES=TRAIN_CLASSES)  
            giou_loss += loss_items[0]  
            conf_loss += loss_items[1]  
            prob_loss += loss_items[2]  
  
        total_loss = giou_loss + conf_loss + prob_loss  
  
    return giou_loss.numpy(), conf_loss.numpy(), prob_loss.numpy(), total_loss.numpy()
```



# TensorFlow\_2.x\_YOLOv3

```
def compute_loss(pred, conv, label, bboxes, i=0, CLASSES=YOLO_COCO_CLASSES):
    NUM_CLASS = len(read_class_names(CLASSES))
    conv_shape = tf.shape(conv)
    batch_size = conv_shape[0]
    output_size = conv_shape[1]
    input_size = STRIDES[i] * output_size
    conv = tf.reshape(conv, (batch_size, output_size, output_size, 3, 5 + NUM_CLASS))

    conv_raw_conf = conv[:, :, :, :, 4:5]
    conv_raw_prob = conv[:, :, :, :, 5:]

    pred_xywh = pred[:, :, :, :, 0:4]
    pred_conf = pred[:, :, :, :, 4:5]

    label_xywh = label[:, :, :, :, 0:4]
    respond_bbox = label[:, :, :, :, 4:5]
    label_prob = label[:, :, :, :, 5:]

    giou = tf.expand_dims(bbox_giou(pred_xywh, label_xywh), axis=-1)
    input_size = tf.cast(input_size, tf.float32)

    bbox_loss_scale = 2.0 - 1.0 * label_xywh[:, :, :, :, 2:3] * label_xywh[:, :, :, :, 3:4] / (input_size ** 2)
    giou_loss = respond_bbox * bbox_loss_scale * (1 - giou)

    iou = bbox_iou(pred_xywh[:, :, :, :, np.newaxis, :], bboxes[:, np.newaxis, np.newaxis, np.newaxis, :, :])
    # Find the value of IoU with the real box The largest prediction box
    max_iou = tf.expand_dims(tf.reduce_max(iou, axis=-1), axis=-1)

    # If the largest iou is less than the threshold, it is considered that the prediction box contains no objects, then the background box
    respond_bgd = (1.0 - respond_bbox) * tf.cast( max_iou < YOLO_IOU_LOSS_THRESH, tf.float32 )

    conf_focal = tf.pow(respond_bbox - pred_conf, 2)

    # Calculate the loss of confidence
    # we hope that if the grid contains objects, then the network output prediction box has a confidence of 1 and 0 when there is no object.
    conf_loss = conf_focal * (
        respond_bbox * tf.nn.sigmoid_cross_entropy_with_logits(labels=respond_bbox, logits=conv_raw_conf)
        +
        respond_bgd * tf.nn.sigmoid_cross_entropy_with_logits(labels=respond_bbox, logits=conv_raw_conf)
    )

    prob_loss = respond_bbox * tf.nn.sigmoid_cross_entropy_with_logits(labels=label_prob, logits=conv_raw_prob)

    giou_loss = tf.reduce_mean(tf.reduce_sum(giou_loss, axis=[1,2,3,4]))
    conf_loss = tf.reduce_mean(tf.reduce_sum(conf_loss, axis=[1,2,3,4]))
    prob_loss = tf.reduce_mean(tf.reduce_sum(prob_loss, axis=[1,2,3,4]))

    return giou_loss, conf_loss, prob_loss
```


- train.py

⇒ main()

⇒ train\_step() v

⇒ compute\_loss() v

# GIoU loss란?







# GIoU loss란?

Main Idea | Loss IoU


Intersection over Union (IoU)



$\text{IoU} = 1$



$\text{IoU} = 0.5$



$\text{IoU} = 0$

Loss =  $1 - \text{IoU}$

0


0.5

1

- IoU를 Loss로 사용하려면  $1 - \text{IoU}$ 를 사용하여 두 박스가 잘 겹칠수록 0에 가까워지도록 한다.  
-> But, 세번째 경우에서 문제 발생! 두 박스가 겹치지 않을 경우는?

# GIoU loss란?

Generalized Intersection over Union (GIoU)




- GIoU를 Loss로 사용할 때에는 Loss = 1 - GIoU 형태로 사용하며 Loss의 최대값은 2, 최소값은 0이 되도록 설정



# TensorFlow\_2.x\_YOLOv3

```
1 !python detect_mnist.py
```

```
2023-02-07 18:47:40.143938: W tensorflow/core/common_runtime/gpu/gpu_bfc_allocator.cc:42]
```



**감사합니다 :)**

**2023.02.08**